

Kernel Methods and Nonlinear Classification


Piyush Rai

CS5350/6350: Machine Learning

September 15, 2011

Kernel Methods: Motivation

- Often we want to **capture nonlinear patterns** in the data
 - Nonlinear Regression: Input-output relationship may not be linear
 - Nonlinear Classification: Classes may not be separable by a linear boundary
- Linear models (e.g., linear regression, linear SVM) are not just rich enough




- **Kernels:** Make linear models work in nonlinear settings
 - By **mapping data to higher dimensions** where it exhibits linear patterns
 - Apply the linear model in the new input space
 - Mapping \equiv changing the feature representation
- **Note:** Such mappings can be expensive to compute in general
 - Kernels give such mappings **for (almost) free**
 - In most cases, the mappings need not be even computed
 - .. using the **Kernel Trick!**

Classifying non-linearly separable data

- Consider this binary classification problem



- Each example represented by a **single feature** x
- No linear separator exists for this data
- Now map each example as $x \rightarrow \{x, x^2\}$
 - Each example now has **two features** (“derived” from the old representation)
- Data now becomes linearly separable in the new representation




- Linear in the new representation \equiv nonlinear in the old representation




Classifying non-linearly separable data

- Let's look at another example:



- Each example defined by a **two features** $\mathbf{x} = \{x_1, x_2\}$
- No linear separator exists for this data
- Now map each example as $\mathbf{x} = \{x_1, x_2\} \rightarrow \mathbf{z} = \{x_1^2, \sqrt{2}x_1x_2, x_2^2\}$
 - Each example now has **three features** ("derived" from the old representation)
- Data now becomes linearly separable in the new representation



Feature Mapping

- Consider the following mapping ϕ for an example $\mathbf{x} = \{x_1, \dots, x_D\}$

$$\phi : \mathbf{x} \rightarrow \{x_1^2, x_2^2, \dots, x_D^2, x_1x_2, x_1x_3, \dots, x_1x_D, \dots, x_{D-1}x_D\}$$

- It's an example of a quadratic mapping
 - Each new feature uses a pair of the original features
- **Problem:** Mapping usually leads to the number of features blow up!
 - Computing the mapping itself can be inefficient in such cases
 - Moreover, *using* the mapped representation could be inefficient too
 - e.g., imagine computing the similarity between two examples: $\phi(\mathbf{x})^\top \phi(\mathbf{z})$
- Thankfully, Kernels help us avoid both these issues!
 - The mapping doesn't have to be explicitly computed
 - Computations with the mapped features remain efficient

Kernels as High Dimensional Feature Mapping

- Consider two examples $\mathbf{x} = \{x_1, x_2\}$ and $\mathbf{z} = \{z_1, z_2\}$
- Let's assume we are given a function k (kernel) that takes as inputs \mathbf{x} and \mathbf{z}

$$\begin{aligned}k(\mathbf{x}, \mathbf{z}) &= (\mathbf{x}^\top \mathbf{z})^2 \\&= (x_1 z_1 + x_2 z_2)^2 \\&= x_1^2 z_1^2 + x_2^2 z_2^2 + 2x_1 x_2 z_1 z_2 \\&= (x_1^2, \sqrt{2}x_1 x_2, x_2^2)^\top (z_1^2, \sqrt{2}z_1 z_2, z_2^2) \\&= \phi(\mathbf{x})^\top \phi(\mathbf{z})\end{aligned}$$

- The above k **implicitly** defines a mapping ϕ to a higher dimensional space

$$\phi(\mathbf{x}) = \{x_1^2, \sqrt{2}x_1 x_2, x_2^2\}$$

- Note that we didn't have to define/compute this mapping
- Simply defining the kernel a certain way gives a higher dim. mapping ϕ
- Moreover the kernel $k(\mathbf{x}, \mathbf{z})$ also computes the dot product $\phi(\mathbf{x})^\top \phi(\mathbf{z})$
 - $\phi(\mathbf{x})^\top \phi(\mathbf{z})$ would otherwise be much more expensive to compute explicitly
- All kernel functions have these properties

Kernels: Formally Defined

- Recall: Each kernel k has an associated feature mapping ϕ
- ϕ takes input $\mathbf{x} \in \mathcal{X}$ (input space) and maps it to \mathcal{F} (“feature space”)
- Kernel $k(\mathbf{x}, \mathbf{z})$ takes two inputs and gives their **similarity** in \mathcal{F} space

$$\phi : \mathcal{X} \rightarrow \mathcal{F}$$

$$k : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}, \quad k(\mathbf{x}, \mathbf{z}) = \phi(\mathbf{x})^\top \phi(\mathbf{z})$$

- \mathcal{F} needs to be a *vector space* with a *dot product* defined on it
 - Also called a **Hilbert Space**
- Can just *any* function be used as a kernel function?
 - No. It must satisfy **Mercer's Condition**

Mercer's Condition

- For k to be a kernel function
 - There must exist a Hilbert Space \mathcal{F} for which k defines a dot product
 - The above is true if K is a **positive definite function**

$$\int d\mathbf{x} \int d\mathbf{z} f(\mathbf{x}) k(\mathbf{x}, \mathbf{z}) f(\mathbf{z}) > 0 \quad (\forall f \in L_2)$$

- This is Mercer's Condition
- Let k_1, k_2 be two kernel functions then the following are as well:
 - $k(\mathbf{x}, \mathbf{z}) = k_1(\mathbf{x}, \mathbf{z}) + k_2(\mathbf{x}, \mathbf{z})$: direct sum
 - $k(\mathbf{x}, \mathbf{z}) = \alpha k_1(\mathbf{x}, \mathbf{z})$: scalar product
 - $k(\mathbf{x}, \mathbf{z}) = k_1(\mathbf{x}, \mathbf{z}) k_2(\mathbf{x}, \mathbf{z})$: direct product
 - Kernels can also be constructed by composing these rules

The Kernel Matrix

- The kernel function k also defines the Kernel Matrix \mathbf{K} over the data
- Given N examples $\{\mathbf{x}_1, \dots, \mathbf{x}_N\}$, the (i, j) -th entry of \mathbf{K} is defined as:

$$K_{ij} = k(\mathbf{x}_i, \mathbf{x}_j) = \phi(\mathbf{x}_i)^\top \phi(\mathbf{x}_j)$$

- K_{ij} : Similarity between the i -th and j -th example in the feature space \mathcal{F}
- \mathbf{K} : $N \times N$ matrix of pairwise similarities between examples in \mathcal{F} space
- \mathbf{K} is a symmetric matrix
- \mathbf{K} is a **positive definite matrix** (except for a few exceptions)
- For a P.D. matrix: $\mathbf{z}^\top \mathbf{K} \mathbf{z} > 0$, $\forall \mathbf{z} \in \mathbb{R}^N$ (also, all eigenvalues positive)
- The Kernel Matrix \mathbf{K} is also known as the **Gram Matrix**

Some Examples of Kernels

The following are the most popular kernels for real-valued vector inputs

- Linear (trivial) Kernel:

$$k(\mathbf{x}, \mathbf{z}) = \mathbf{x}^\top \mathbf{z} \text{ (mapping function } \phi \text{ is identity - no mapping)}$$

- Quadratic Kernel:

$$k(\mathbf{x}, \mathbf{z}) = (\mathbf{x}^\top \mathbf{z})^2 \quad \text{or} \quad (1 + \mathbf{x}^\top \mathbf{z})^2$$

- Polynomial Kernel (of degree d):

$$k(\mathbf{x}, \mathbf{z}) = (\mathbf{x}^\top \mathbf{z})^d \quad \text{or} \quad (1 + \mathbf{x}^\top \mathbf{z})^d$$

- Radial Basis Function (RBF) Kernel:

$$k(\mathbf{x}, \mathbf{z}) = \exp[-\gamma \|\mathbf{x} - \mathbf{z}\|^2]$$

- γ is a hyperparameter (also called the **kernel bandwidth**)
- The RBF kernel corresponds to an **infinite dimensional** feature space \mathcal{F} (i.e., you can't actually write down the vector $\phi(\mathbf{x})$)

Note: Kernel hyperparameters (e.g., d , γ) chosen via cross-validation

Using Kernels

- Kernels can turn a linear model into a nonlinear one
- Recall: Kernel $k(\mathbf{x}, \mathbf{z})$ represents a dot product in some high dimensional feature space \mathcal{F}
- Any learning algorithm **in which examples only appear as dot products** ($\mathbf{x}_i^\top \mathbf{x}_j$) can be kernelized (i.e., non-linearized)
 - .. by replacing the $\mathbf{x}_i^\top \mathbf{x}_j$ terms by $\phi(\mathbf{x}_i)^\top \phi(\mathbf{x}_j) = k(\mathbf{x}_i, \mathbf{x}_j)$
- Most learning algorithms are like that
 - Perceptron, SVM, linear regression, etc.
 - Many of the unsupervised learning algorithms too can be kernelized (e.g., K -means clustering, Principal Component Analysis, etc.)

Kernelized SVM Training

- Recall the SVM dual Lagrangian:

$$\begin{aligned} \text{Maximize } L_D(\mathbf{w}, b, \xi, \alpha, \beta) &= \sum_{n=1}^N \alpha_n - \frac{1}{2} \sum_{m,n=1}^N \alpha_m \alpha_n y_m y_n (\mathbf{x}_m^T \mathbf{x}_n) \\ \text{subject to } \sum_{n=1}^N \alpha_n y_n &= 0, \quad 0 \leq \alpha_n \leq C; \quad n = 1, \dots, N \end{aligned}$$

- Replacing $\mathbf{x}_m^T \mathbf{x}_n$ by $\phi(\mathbf{x}_m)^T \phi(\mathbf{x}_n) = k(\mathbf{x}_m, \mathbf{x}_n) = K_{mn}$, where $k(.,.)$ is some suitable kernel function

$$\begin{aligned} \text{Maximize } L_D(\mathbf{w}, b, \xi, \alpha, \beta) &= \sum_{n=1}^N \alpha_n - \frac{1}{2} \sum_{m,n=1}^N \alpha_m \alpha_n y_m y_n K_{mn} \\ \text{subject to } \sum_{n=1}^N \alpha_n y_n &= 0, \quad 0 \leq \alpha_n \leq C; \quad n = 1, \dots, N \end{aligned}$$

- SVM now learns a linear separator in the kernel defined feature space \mathcal{F}
 - This corresponds to a non-linear separator in the original space \mathcal{X}

Kernelized SVM Prediction

- Prediction for a test example \mathbf{x} (assume $b = 0$)

$$y = \text{sign}(\mathbf{w}^\top \mathbf{x}) = \text{sign}\left(\sum_{n \in SV} \alpha_n y_n \mathbf{x}_n^\top \mathbf{x}\right)$$

- SV is the set of support vectors (i.e., examples for which $\alpha_n > 0$)
- Replacing each example with its feature mapped representation ($\mathbf{x} \rightarrow \phi(\mathbf{x})$)


$$y = \text{sign}\left(\sum_{n \in SV} \alpha_n y_n \phi(\mathbf{x}_n)^\top \phi(\mathbf{x})\right) = \text{sign}\left(\sum_{n \in SV} \alpha_n y_n k(\mathbf{x}_n, \mathbf{x})\right)$$

- The weight vector for the kernelized case can be expressed as:

$$\mathbf{w} = \sum_{n \in SV} \alpha_n y_n \phi(\mathbf{x}_n) = \sum_{n \in SV} \alpha_n y_n k(\mathbf{x}_n, \cdot)$$

- Important:** Kernelized SVM needs the support vectors at the test time (except when you can write $\phi(\mathbf{x}_n)$ as an explicit, reasonably-sized vector)
 - In the unkernelized version $\mathbf{w} = \sum_{n \in SV} \alpha_n y_n \mathbf{x}_n$ can be computed and stored as a $D \times 1$ vector, so the support vectors need not be stored

SVM with an RBF Kernel



- The learned decision boundary in the original space is nonlinear

Kernels: concluding notes

- Kernels give a modular way to learn nonlinear patterns using linear models
 - All you need to do is replace the inner products with the kernel
- All the computations remain as efficient as in the original space
- Choice of the kernel is an important factor
- Many kernels are tailor-made for specific types of data
 - Strings (string kernels): DNA matching, text classification, etc.
 - Trees (tree kernels): Comparing parse trees of phrases/sentences
- Kernels can even be **learned from the data** (**hot research topic!**)
 - Kernel learning means **learning the similarities** between examples (instead of using some pre-defined notion of similarity)
- **A question worth thinking about:** Wouldn't mapping the data to higher dimensional space cause my classifier (say SVM) to overfit?
 - The answer lies in the concepts of **large margins** and **generalization**

Next class..

- Intro to probabilistic methods for supervised learning
 - Linear Regression (probabilistic version)
 - Logistic Regression